

1. What is a Markov Decision Process (MDP) and what are its components (states, actions, transition, reward) ?
2. How do you prevent overfitting in deep neural networks ?
3. What is the bias-variance tradeoff in machine learning, and how does it affect model performance ?
4. Explain the concept of a probability density function (PDF) vs. a cumulative distribution function (CDF) .
5. What is a prior distribution, and how does it affect Bayesian inference ?
6. How does a neural network learn? Describe the backpropagation algorithm and how gradients are used .
7. What is an ML feature store, and why might it be useful ?
8. How do ensemble methods like bagging and boosting work, and when might you use them ?
9. Explain Continuous Integration/Continuous Delivery (CI/CD) in MLOps. How do they apply to ML ?
10. What is reward shaping and why might it be used ?
11. What is dropout and why is it used? What is batch normalization and how does it help training ?
12. What are regularization techniques (e.g., L1, L2), and how do they prevent overfitting ?
13. How would you handle missing data in a probabilistic framework ?

14. How do you evaluate an NLP model such as a translation or summarization system (mention BLEU, ROUGE, etc.) ?
15. What are embeddings in deep learning (e.g., word embeddings, image embeddings), and why are they useful ?
16. Explain the vanishing/exploding gradient problem and how it affects training deep networks .
17. What is cross-validation, and why is it used in model evaluation ?
18. What is the difference between generative and discriminative models? Give examples of each .
19. What is word sense disambiguation and how might an NLP system handle it ?
20. What is algorithmic bias, and how can it arise in AI systems ?
21. What is an activation function? Compare common ones (ReLU, sigmoid, tanh) and explain why deep networks often use ReLU .
22. How would you mitigate bias in a dataset or model ?
23. What are some common image preprocessing techniques before feeding data into a neural network ?
24. Explain classification metrics such as precision, recall, and F1 score, and when you would prefer one over the others .
25. Describe how gradient descent variants (SGD, Adam, RMSprop) differ and when you might choose each .
26. What is the role of learning rate in training a neural network? How do you choose or adapt it ?

27. Describe RNNS, LSTMs, and GRUS. Why are LSTMs/GRUs used for sequential data ?
28. What is a Bayesian network (belief network)? How does it represent dependencies between variables ?
29. How do you evaluate a model's performance? Describe metrics for classification (e.g., accuracy, precision, recall, ROC-AUC) and regression (e.g., MSE,  $R^2$ ) .
30. Explain underfitting and overfitting. How can you detect and address each ?
31. What is the exploration vs exploitation dilemma in RL, and how can it be addressed (e.g.,  $\epsilon$ -greedy, UCB) ?
32. What is transfer learning in deep learning, and when would you use it ?
33. Explain the difference between bag-of-words and word embeddings (e.g., word2vec, GloVe). Why are embeddings powerful ?
34. Explain actor-critic algorithms at a high level .
35. How does gradient descent work? What are some challenges associated with it (e.g., local minima, learning rate choice) ?
36. What is a graphical model? How does it help in probabilistic reasoning ?
37. What are some challenges of RL (e.g., sample inefficiency, sparse rewards) ?
38. How does an object detection model (e.g., YOLO, Faster R-CNN) differ from an image classification model ?
39. What is a language model? How do n-gram models differ from modern transformer-based models ?

40. What is the No Free Lunch Theorem in machine learning, and what does it imply for model selection ?

41. What is explainability or interpretability in AI, and why is it important ?

42. How does the TF-IDF statistic work, and when might you use it ?

43. What is model governance, and why is it important in ML workflows ?

44. Explain different notions of fairness in machine learning (e.g., equalized odds, demographic parity) and potential trade-offs .

45. What strategies exist to make AI systems more transparent or explainable (e.g., model-agnostic methods like LIME/SHAP) ?

46. What is an autoencoder? How can autoencoders be used for tasks like dimensionality reduction or denoising ?

47. What is semantic segmentation versus instance segmentation ?

48. Why are GPU/TPU accelerators especially beneficial for computer vision tasks ?

49. What are Intersection over Union (IoU) and mean Average Precision (mAP), and how are they used in object detection evaluation ?

50. Describe the assumptions of linear regression and how you would check them .

51. What is data augmentation in computer vision, and why is it helpful ?

52. What is tokenization in NLP, and why is it important for text processing ?

53. What does the term "ethical AI" mean to you, and what principles are important (e.g., transparency, accountability, justice) ?

54. Explain feature selection and feature engineering. What methods can be used for selecting features ?
55. How can you automate retraining and redeployment of models when new data arrives ?
56. Explain how containerization (e.g., Docker) and orchestration (e.g., Kubernetes) are used in ML deployment .
57. Explain how sampling methods (e.g., Monte Carlo, Gibbs sampling) are used in probabilistic models .
58. Explain the difference between on-policy and off-policy learning (e.g., SARSA vs Q-learning) .
59. Define ROC curve and AUC. What do they tell you about a binary classifier ?
60. What is a convolution operation in a CNN, and how do filters (kernels) work ?
61. What is the curse of dimensionality, and how can it impact machine learning algorithms ?
62. Explain the attention mechanism and how it is used in Transformer architectures (e.g., BERT, GPT) .
63. What is model drift (data drift or concept drift), and how can you detect and mitigate it ?
64. Describe a Gaussian Mixture Model and its use cases .
65. What is the Bellman equation for a value function ?
66. How does temporal difference (TD) learning work ?

67. Describe Generative Adversarial Networks (GANs). How do the generator and discriminator work together ?
68. Explain a Markov chain and a Hidden Markov Model (HMM) in broad terms .
69. What is the role of a feature extractor vs a classifier in CNN-based pipelines ?
70. How do non-max suppression algorithms work for selecting bounding boxes in detection ?
71. Why is monitoring of deployed ML models important, and which metrics would you monitor ?
72. How do transformers (like BERT or GPT) differ from RNNs in processing text ?
73. How do you version control datasets and models in an ML project ?
74. How does Deep Q-Network (DQN) differ from classical Q-learning ?
75. What is informed consent in data collection, and why is it important for AI ethics ?
76. What are the main steps of the ML development lifecycle (data, training, deployment, monitoring) ?
77. Explain pooling layers (max pooling vs average pooling) and their effect on feature maps .
78. What are some common NLP preprocessing steps (e.g., stop-word removal, stemming, lemmatization) and why are they used ?
79. How do privacy regulations like GDPR impact AI data collection and model building ?

80. Describe the Expectation-Maximization (EM) algorithm and when you would use it .
81. Describe policy gradient methods and how they differ from value-based methods .
82. What is the difference between a generative model and a discriminative model? Give examples .
83. Explain transfer learning in the context of CV (e.g., using pre-trained CNNs on ImageNet for new tasks) .
84. Describe canary deployments or A/B testing for ML models in production .
85. What are function approximators (e.g., neural networks) used for in RL ?
86. State Bayes' theorem and explain each term (prior, likelihood, posterior). Examples.
87. What is maximum likelihood estimation (MLE)? How does it differ from maximum a posteriori (MAP) estimation ?
88. What is the difference between batch inference and real-time inference, and when would you use each ?
89. Explain self-supervised learning in NLP (e.g., masked language modeling) and why it's useful .
90. Describe a popular CNN architecture (e.g., ResNet, VGG) and its key ideas (like residual connections) .
91. Discuss the trade-off between model performance and fairness. Can they conflict ?

92. Give an example of an AI application that raised ethical concerns, and how those concerns were addressed .
93. What is PCA (Principal Component Analysis) and when would you use it ?
94. How might an AI model inadvertently cause harm (e.g., reinforcement of stereotypes, safety issues) ?
95. Compare Convolutional Neural Networks (CNNs) and Recurrent Neural Networks (RNNs) in terms of architecture and applications .
96. How can AI systems discriminate against protected groups, and how would you test for it ?
97. Explain discounted reward and the role of the discount factor in RL .
98. What is the attention mechanism, and how has it improved NLP models ?
99. What are some challenges unique to deploying ML models compared to traditional software ?
100. Explain the architecture of a sequence-to-sequence (seq2seq) model with attention. What tasks does it solve ?